

ORION WELLER

Baltimore Maryland, 21218

<http://orionweller.github.io> ♦ wellerorion@gmail.com

EDUCATION

Johns Hopkins University Ph.D. in Computer Science Center for Language and Speech Processing Advisors: Benjamin Van Durme and Dawn Lawrie	08/2021 – 05/2026 (expected) GPA: 4.0/4.0
Brigham Young University Bachelor of Science, Summa Cum Laude (Top 1%) Majors: Computer Science, Statistics; Minor in Mathematics	08/2017 – 04/2021 GPA: 4.0/4.0

AWARDS

- NSF Graduate Research Fellowship, 2022
- DoD NDSEG Fellowship (declined), 2022
- CRA Outstanding Undergraduate Researcher in North America Award, 2021
- Fulbright Fellowship Semi-Finalist, 2021
- Goldwater Scholarship, 2020

INDUSTRY EXPERIENCE

Samaya AI <i>PhD Intern</i> , with Jack Hessel, Ashwin Paranjape, and Yuhao Zhang	Summer 2024
Allen Institute for Artificial Intelligence (AI2) <i>Research Intern</i> , with Luca Soldaini, Kyle Lo, and Arman Cohan	Summer 2023
Apple <i>Machine Learning Intern in AI/ML</i> , with Matthias Sperber	Summers 2020, 2021
Allen Institute for Artificial Intelligence (AI2) <i>Research Intern</i> , with Matt Gardner and Matt Peters	Winter 2020
Qualtrics <i>Software Engineer Intern</i>	Summer 2019
Adobe <i>Student Researcher</i> , with Ilya Reznik and Chris Challis	Fall 2018

PUBLICATIONS

- [1] **Orion Weller**, Benjamin Van Durme, Dawn Lawrie, Ashwin Paranjape, Yuhao Zhang, Jack Hessel, “Promptriever: Instruction-Trained Retrievers Can Be Prompted Like Language Models,” in *ArXiv*, 2024.
- [2] **Orion Weller**, Benjamin Chang, Sean MacAvaney, Kyle Lo, Arman Cohan, Benjamin Van Durme, Dawn Lawrie, Luca Soldaini, “FollowIR: Evaluating and Teaching Information Retrieval Models to Follow Instructions,” in *ArXiv*, 2024.
- [3] Jeffrey Cheng, Marc Marone, **Orion Weller**, Dawn Lawrie, Daniel Khashabi, Benjamin Van Durme, “Dated Data: Tracing Knowledge Cutoffs in Large Language Models,” in *Conference on Language Models (CoLM)*, 2024 **Outstanding Paper Award (top 4/300)**.

- [4] Nathaniel Weir, Muhammad Khalifa, Linlu Qiu, **Orion Weller**, Peter Clark, “Learning to Reason via Program Generation, Emulation, and Search,” in *NeurIPS*, 2024.
- [5] Brennon Brimhall, **Orion Weller**, Matthew Green, Ian Miers, “A Waterlog for Detecting and Tracing Synthetic Text from Large Language Models,” in *Cryptology ePrint Archive*, 2024.
- [6] Abe Bohan Hou, **Orion Weller**, Guanghui Qin, Eugene Yang, Dawn Lawrie, Nils Holzenberger, Andrew Blair-Stanek, Benjamin Van Durme, “CLERC: A Dataset for Legal Case Retrieval and Retrieval-Augmented Analysis Generation,” in *ArXiv*, 2024.
- [7] James Mayfield, Eugene Yang, Dawn Lawrie, Sean MacAvaney, Paul McNamee, Douglas W Oard, Luca Soldaini, Ian Soboroff, **Orion Weller**, Efsun Kayi, Kate Sanders, Marc Mason, Noah Hibbler, “On the Evaluation of Machine-Generated Reports,” in *SIGIR*, 2024.
- [8] Dongwei Jiang, Jingyu Zhang, **Orion Weller**, Nathaniel Weir, Benjamin Van Durme, Daniel Khashabi, “Self-[In]correct: LLMs Struggle with Refining Self-Generated Responses,” 2024.
- [9] Nathaniel Weir, Kate Sanders, **Orion Weller**, Shreya Sharma, Dongwei Jiang, Zhengping Zhang, Bhavana Dalvi Mishra, Oyvind Tafjord, Peter Jansen, Peter Clark, “Enhancing Systematic Decompositional Natural Language Inference Using Informal Logic,” in *EMNLP*, 2024.
- [10] **Orion Weller**, Dawn Lawrie, Benjamin Van Durme, “NevIR: Negation in Neural Information Retrieval,” in *EACL*, 2024 **Selected for Oral**.
- [11] **Orion Weller**, Kyle Lo, David Wadden, Dawn Lawrie, Benjamin Van Durme, Arman Cohan, Luca Soldaini, “When do Generative Query and Document Expansions Fail? A Comprehensive Study Across Methods, Retrievers, and Datasets,” in *Findings of EACL*, 2024.
- [12] **Orion Weller**, Marc Marone, Nathaniel Weir, Dawn Lawrie, Daniel Khashabi, Benjamin Van Durme, ““According to...” Prompting Language Models Improves Quoting from Pre-Training Data,” in *EACL*, 2024.
- [13] **Orion Weller**, Aleem Khan, Nathaniel Weir, Dawn J Lawrie, Benjamin Van Durme, “Defending Against Disinformation Attacks in Open-Domain Question Answering,” in *EACL*, 2024 **Selected for Oral**.
- [14] James Mayfield, Eugene Yang, Dawn Lawrie, Samuel Barham, **Orion Weller**, Marc Mason, Suraj Nair, Scott Miller, “Synthetic Cross-language Information Retrieval Training Data,” in *ArXiv*, 2023.
- [15] Samuel Barham, **Orion Weller**, Michelle Yuan, Kenton Murray, Mahsa Yarmohammadi, Zhengping Jiang, Siddharth Vashishtha, Alexander Martin, Anqi Liu, Aaron Steven White, “MegaWika: Millions of Reports and their Sources across 50 Diverse Languages,” in *ArXiv*, 2023.
- [16] Kangda Wei, Dawn J Lawrie, Benjamin Van Durme, Yunmo Chen, **Orion Weller**, “When Do Decompositions Help for Machine Reading?” In *EMNLP*, 2023.
- [17] **Orion Weller**, Marc Marone, Vladimir Braverman, Dawn Lawrie, Benjamin Van Durme, “Pre-trained Models for Multilingual Federated Learning,” in *NAACL*, 2022.
- [18] **Orion Weller**, Kevin Seppi, Matt Gardner, “When to Use Multi-Task Learning vs Intermediate Fine-Tuning for Pre-Trained Encoder Transfer Learning,” in *ACL*, 2022. **Selected for Oral**.
- [19] **Orion Weller**, Matthias Sperber, Telmo Pires, Hendra Setiawan, Christian Gollan, Dominic Telaar, Matthias Paulik, “Speech Translation for Code-Switched Speech,” in *Findings of the ACL*, 2022.
- [20] Wilson Fearn, **Orion Weller**, Kevin Seppi, “Exploring the Relationship Between Algorithm Performance, Vocabulary, and Run-Time in Text Classification,” in *NAACL*, 2021.

- [21] **Orion Weller**, Matthias Sperber, Christian Gollan, Joris Kluivers, “Streaming Joint Speech Translation and Transcription,” in *EACL*, 2021.
- [22] **Orion Weller**, Luke Sagers, Carl Hanson, Quinn Snell, Michael Barnes, Shannon Tass, “Predicting Suicidal Thoughts and Behavior among Adolescents using the Risk and Protective Factor Framework: a Large-Scale Machine Learning Approach,” in *Plos One*, 2021.
- [23] **Orion Weller**, Nicholas Lourie, Matt Gardner, Matthew Peters, “Learning from Task Descriptions,” in *EMNLP*, 2020.
- [24] **Orion Weller**, Jordan Hildebrandt, Ilya Reznik, Christopher Challis, E. Shannon Tass, Quinn Snell, Kevin Seppi, “You Don’t Have Time to Read This: an Exploration of Document Level Reading Time Prediction,” in *ACL*, 2020.
- [25] **Orion Weller**, Nancy Fulda, Kevin Seppi, “Can Humor Prediction Datasets be used for Humor Generation? Humorous Headline Generation via Style Transfer,” in *FigLang @ ACL*, 2020.
- [26] **Orion Weller**, Kevin Seppi, “The rJokes Dataset: a Large Scale Humor Collection,” in *LREC*, 2020.
- [27] **Orion Weller**, Kevin Seppi, “Humor Detection: a Transformer Gets the Last Laugh,” in *EMNLP*, 2019.

SKILLS

Languages	English (native), Portuguese (intermediate), French (basic)
Programming Languages	Python, C++, R, Javascript, Golang, Java, OCaml, SQL
Frameworks & Tools	PyTorch, Linux, Mechanical Turk, Django, VueJS

SERVICE

Workshop Organizer

- Scholarly Document Processing (SDP) Workshop at ACL, 2024
- Mid-Atlantic Student Colloquium on Speech, Language and Learning, 2022-2024

Teaching

- EN.500.111: Reasoning with ChatGPT, Johns Hopkins University, Fall 2023
- After School Coding Instructor for Elementary School with Code in the Schools, 2021-2022

Program Committee Member

- 2024: ARR, NeurIPS, CoLM, LLM4Eval
- 2023: ACL, EMNLP, ARR, NeurIPS
- 2022: ACL Rolling Review (ARR)
- 2021: ACL Rolling Review, ACL, EMNLP, NAACL (secondary)
- 2020: Meta-Learning@NeurIPS